

## **Machine Learning en utilisant Sklearn et Pandas**

### **1. Introduction**

Dans ce tutorial, nous allons implémenter des scripts Python en utilisant deux packages Sklearn et Pandas pour effectuer l'apprentissage automatique.

En particulier, nous allons voir des exemples de 3 tâches de base de l'apprentissage automatique : **Clustering**, **Classification** et **Regression**.

## 2. Clustering


```
from sklearn.cluster import KMeans
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from matplotlib import pyplot as plt
%matplotlib inline
```

```
df = pd.read_csv("income.csv")
df.head()
```

	Name	Age	Income(\$)
0	Rob	27	70000
1	Michael	29	90000
2	Mohan	29	61000
3	Ismail	28	60000
4	Kory	42	150000

```
plt.scatter(df.Age, df['Income($)'])
plt.xlabel('Age')
plt.ylabel('Income($)')
```

```
Text(0, 0.5, 'Income($)')
```



```
km = KMeans(n_clusters=3)
y_predicted = km.fit_predict(df[['Age', 'Income($)']])
y_predicted

array([0, 0, 2, 2, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 0, 0, 2])
```

```
df['cluster']=y_predicted
df.head()
```


	Name	Age	Income(\$)	cluster
0	Rob	27	70000	0
1	Michael	29	90000	0
2	Mohan	29	61000	2
3	Ismail	28	60000	2
4	Kory	42	150000	1

```
km.cluster_centers_
```

```
array([[3.40000000e+01, 8.05000000e+04],
       [3.82857143e+01, 1.50000000e+05],
       [3.29090909e+01, 5.61363636e+04]])
```

```
df1 = df[df.cluster==0]
df2 = df[df.cluster==1]
df3 = df[df.cluster==2]
plt.scatter(df1.Age,df1['Income($)'],color='green')
plt.scatter(df2.Age,df2['Income($)'],color='red')
plt.scatter(df3.Age,df3['Income($)'],color='black')
plt.scatter(km.cluster_centers_[:,0],km.cluster_centers_[:,1],color='purple',marker='*',label='centroid')
plt.xlabel('Age')
plt.ylabel('Income ($)')
plt.legend()
```

<matplotlib.legend.Legend at 0x279134e1860>



## Clustering avec pré-traitement :

```
scaler = MinMaxScaler()

scaler.fit(df[['Income($)']])
df['Income($)'] = scaler.transform(df[['Income($)']])


scaler.fit(df[['Age']])
df['Age'] = scaler.transform(df[['Age']])
```

```
df.head()
```

	Name	Age	Income(\$)	cluster
0	Rob	0.058824	0.213675	1
1	Michael	0.176471	0.384615	1
2	Mohan	0.176471	0.136752	1
3	Ismail	0.117647	0.128205	1
4	Kory	0.941176	0.897436	0

```
plt.scatter(df.Age, df['Income($)'])
```

```
<matplotlib.collections.PathCollection at 0x279136d9cc0>
```



```
km = KMeans(n_clusters=3)
y_predicted = km.fit_predict(df[['Age', 'Income($)']])
y_predicted
```

```
array([1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0])
```

```
df['cluster'] = y_predicted
df.head()
```

	Name	Age	Income(\$)	cluster
0	Rob	0.058824	0.213675	1
1	Michael	0.176471	0.384615	1
2	Mohan	0.176471	0.136752	1
3	Ismail	0.117647	0.128205	1
4	Kory	0.941176	0.897436	0

```
km.cluster_centers_
```


```
array([[0.72268908, 0.8974359 ],
       [0.1372549 , 0.11633428],
       [0.85294118, 0.2022792 ]])
```

```

df1 = df[df.cluster==0]
df2 = df[df.cluster==1]
df3 = df[df.cluster==2]
plt.scatter(df1.Age,df1['Income($)'],color='green')
plt.scatter(df2.Age,df2['Income($)'],color='red')
plt.scatter(df3.Age,df3['Income($)'],color='black')
plt.scatter(km.cluster_centers[:,0],km.cluster_centers[:,1],color='purple',marker='*',label='centroid')
plt.legend()

```

<matplotlib.legend.Legend at 0x27913611e80>



### 3. Classification

```
import pandas as pd
from sklearn.datasets import load_iris
iris = load_iris()
```

```
iris.feature_names
```

```
['sepal length (cm)',
 'sepal width (cm)',
 'petal length (cm)',
 'petal width (cm)']
```

```
iris.target_names
```

```
array(['setosa', 'versicolor', 'virginica'], dtype='<U10')

```

```
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

```
df['target'] = iris.target
df.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
df[df.target==1].head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
50	7.0	3.2	4.7	1.4	1
51	6.4	3.2	4.5	1.5	1
52	6.9	3.1	4.9	1.5	1
53	5.5	2.3	4.0	1.3	1
54	6.5	2.8	4.6	1.5	1

```
df[df.target==2].head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
100	6.3	3.3	6.0	2.5	2
101	5.8	2.7	5.1	1.9	2
102	7.1	3.0	5.9	2.1	2
103	6.3	2.9	5.6	1.8	2
104	6.5	3.0	5.8	2.2	2

```
df['flower_name'] =df.target.apply(lambda x: iris.target_names[x])
df.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	flower_name
0	5.1	3.5	1.4	0.2	0	setosa
1	4.9	3.0	1.4	0.2	0	setosa
2	4.7	3.2	1.3	0.2	0	setosa
3	4.6	3.1	1.5	0.2	0	setosa
4	5.0	3.6	1.4	0.2	0	setosa


```
df[45:55]
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	flower_name
45	4.8	3.0	1.4	0.3	0	setosa
46	5.1	3.8	1.6	0.2	0	setosa
47	4.6	3.2	1.4	0.2	0	setosa
48	5.3	3.7	1.5	0.2	0	setosa
49	5.0	3.3	1.4	0.2	0	setosa
50	7.0	3.2	4.7	1.4	1	versicolor
51	6.4	3.2	4.5	1.5	1	versicolor
52	6.9	3.1	4.9	1.5	1	versicolor
53	5.5	2.3	4.0	1.3	1	versicolor
54	6.5	2.8	4.6	1.5	1	versicolor

### Sepal length vs Sepal Width (Setosa vs Versicolor)

```
plt.xlabel('Sepal Length')
plt.ylabel('Sepal Width')
plt.scatter(df0['sepal length (cm)'], df0['sepal width (cm)'],color="green",marker='+')
plt.scatter(df1['sepal length (cm)'], df1['sepal width (cm)'],color="blue",marker='.')
```


<matplotlib.collections.PathCollection at 0x1ef65816e10>



### Petal length vs Petal Width (Setosa vs Versicolor)

```
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
plt.scatter(df0['petal length (cm)'], df0['petal width (cm)'],color="green",marker='+')
plt.scatter(df1['petal length (cm)'], df1['petal width (cm)'],color="blue",marker='.')
```

<matplotlib.collections.PathCollection at 0x1ef658bbc50>



### Train Using Support Vector Machine (SVM)

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop(['target', 'flower_name'], axis='columns')
y = df.target
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
len(X_train)
```

120

```
len(X_test)
```

30



```
from sklearn.svm import SVC
model = SVC()
```

```
model.fit(X_train, y_train)
```

```
C:\Users\marouane\Anaconda3\lib\site-packages\sklearn\svm\base.py:196:
om 'auto' to 'scale' in version 0.22 to account better for unscaled fea
d this warning.
"avoid this warning.", FutureWarning)
```

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
    kernel='rbf', max_iter=-1, probability=False, random_state=None,
    shrinking=True, tol=0.001, verbose=False)
```

```
model.score(X_test, y_test)
```

```
0.9666666666666667
```

```
model.predict([[4.8,3.0,1.5,0.3]])
```

```
array([0])
```

## 4. Régression linéaire

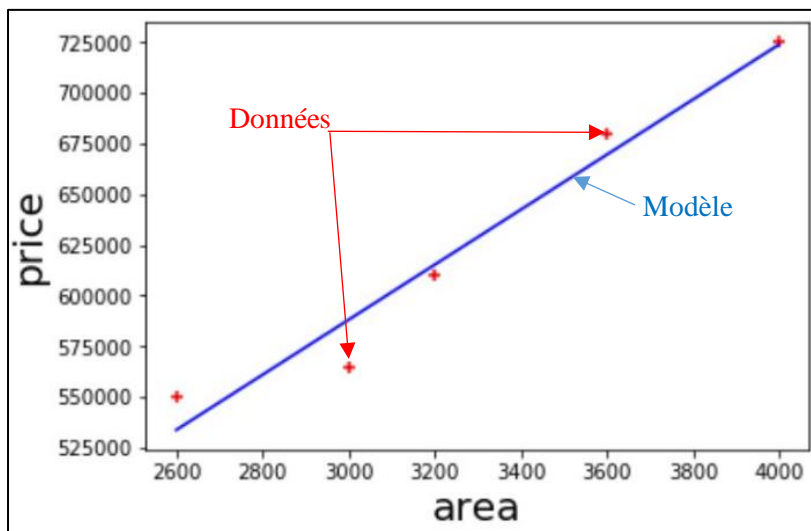
### 4.1. Regression à une seule variable

#### Données


area	price
2600	550000
3000	565000
3200	610000
3600	680000
4000	725000

#### Modèle


**Énoncé du problème:** Etant donné les données ci-dessus, nous allons construire un modèle d'apprentissage automatique capable de prévoir le prix des maisons en fonction de la superficie en m2.



**Solution :**



## Modèle mathématique:



## Implémentation d'apprentissage:


```
import pandas as pd
import numpy as np
from sklearn import linear_model
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('homeprices.csv')
df
```

	area	price
0	2600	550000
1	3000	565000
2	3200	610000
3	3600	680000
4	4000	725000

```
%matplotlib inline
plt.xlabel('area')
plt.ylabel('price')
plt.scatter(df.area, df.price, color='red', marker='+')
```

<matplotlib.collections.PathCollection at 0x1fa581927f0>



```
area = df[['area']]
area
```

	area
0	2600
1	3000
2	3200
3	3600
4	4000

```
price = df.price
price
```

0	550000
1	565000
2	610000
3	680000
4	725000

Name: price, dtype: int64

```
# Create linear regression object
reg = linear_model.LinearRegression()
reg.fit(area,price)
```

LinearRegression(copy\_X=True, fit\_intercept=True, n\_jobs=None, normalize=False)

```
reg.predict([[3000]])
```

array([587979.45205479])

```
reg.coef_
```

array([135.78767123])

```
reg.intercept_
```

180616.43835616432

**Y = m \* X + b (m is coefficient and b is intercept)**

**5000\*135.78767123 + 180616.43835616432**

859554.7945061643

## Implémentation du test :

```
area_df = pd.read_csv("areas.csv")
area_df.head(3)
```

	area
0	1000
1	1500
2	2300

```
p = reg.predict(area_df)
p
```

array([ 316404.10958904, 384297.94520548, 492928.08219178,  
 661304.79452055, 740061.64383562, 799808.21917808,  
 926090.75342466, 650441.78082192, 825607.87671233,  
 492928.08219178, 1402705.47945205, 1348390.4109589 ,  
 1144708.90410959])


```
area_df['prices']=p
area_df
```

	area	prices
0	1000	3.164041e+05
1	1500	3.842979e+05
2	2300	4.929281e+05
3	3540	6.613048e+05
4	4120	7.400616e+05
5	4560	7.998082e+05
6	5490	9.260908e+05
7	3460	6.504418e+05
8	4750	8.256079e+05
9	2300	4.929281e+05
10	9000	1.402705e+06
11	8600	1.348390e+06
12	7100	1.144709e+06

```
area_df.to_csv("prediction.csv")
```

```
%matplotlib inline
plt.xlabel('area', fontsize=20)
plt.ylabel('price', fontsize=20)
plt.scatter(df.area,df.price,color='red',marker='+')
plt.plot(df.area,reg.predict(area),color='blue')
```

```
[<matplotlib.lines.Line2D at 0x27c3c32bc88>]
```



### Exercice :

#### Enoncé :

Prédictez le revenu par habitant du canada en 2020.

Les données sont enregistrées dans le fichier `canada_per_capita_income.csv`.

À partir de là :

1. On construit un modèle de régression
2. On prédit le revenu par habitant des citoyens canadiens en 2020

Réponse :

41288.69409442

## 4.2. Regression à variable multiple

Données

area	bedrooms	age	price
2600	3	20	550000
3000	4	15	565000
3200		18	610000
3600	3	30	595000
4000	5	8	760000
4100	6	8	810000

Modèle mathématique

Dependent variable      Independent variables (**features**)

$$\text{price} = m_1 * \text{area} + m_2 * \text{bedrooms} + m_3 * \text{age} + b$$

Coefficients

$$y = m_1 x_1 + m_2 x_2 + m_3 x_3 + b$$

Implémentation :

```
import pandas as pd
import numpy as np
from sklearn import linear_model
```

```
df = pd.read_csv('homeprices.csv')
df
```

	area	bedrooms	age	price
0	2600	3.0	20	550000
1	3000	4.0	15	565000
2	3200	NaN	18	610000
3	3600	3.0	30	595000
4	4000	5.0	8	760000
5	4100	6.0	8	810000

## Pré-traitement des données :

On va remplacer les valeurs NaN (indéfinies) avec la valeur de médiane de la colonne

```
import math
med_bedrooms = math.floor(df.bedrooms.median())
```

```
df.bedrooms = df.bedrooms.fillna(med_bedrooms)
```

df

	area	bedrooms	age	price
0	2600	3.0	20	550000
1	3000	4.0	15	565000
2	3200	4.0	18	610000
3	3600	3.0	30	595000
4	4000	5.0	8	760000
5	4100	6.0	8	810000

## Apprentissage

```
reg = linear_model.LinearRegression()
reg.fit(df[['area', 'bedrooms', 'age']], df.price)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                  normalize=False)
```

reg.coef\_

```
array([ 112.06244194, 23388.88007794, -3231.71790863])
```

reg.intercept\_

```
221323.00186540408
```

**Test :** trouver le prix d'une maison de superficie 3000 m2 , ayant 3 bedrooms et d'âge 40 ans ?

```
reg.predict([[3000, 3, 40]])
```

```
array([498408.25158031])
```

```
112.06244194*3000 + 23388.88007794*3 + -3231.71790863*40 + 221323.00186540384
```

```
498408.25157402386
```

**Find price of home with 2500 sqr ft area, 4 bedrooms, 5 year old**

```
reg.predict([[2500, 4, 5]])
```

```
array([578876.03748933])
```

## Exercice :

### Enoncé :

Les données sont enregistrées dans le fichier **hiring.csv**.

Ce fichier contient des statistiques d'embauche pour une entreprise telles que l'expérience du candidat, sa note au test écrit et sa note pour l'entretien personnel.

Sur la base de ces 3 facteurs, les ressources humaines décideront du salaire.

Compte tenu de ces données, on doit créer un modèle d'apprentissage automatique pour le service des ressources humaines, qui peut l'aider à déterminer les salaires des futurs candidats.

Travail demandé : Prédire les salaires pour les candidats suivants :

<b>2 yr experience, 9 test score, 6 interview score</b>
<b>12 yr experience, 10 test score, 10 interview score</b>

Réponse :

53713.86

93747.79